

Day6 Delivery and Verification

Deng Jinzhu

Frozen numbers

Baseline TCO: D0 USD 365,811, D1 USD 348,040, D2 USD 356,665. D1 minus D0: TCO -17,771 and P90 -20,908.

D1 decomposition: procurement 318,240 + holding 3,590 + shortage loss 2 + renewal 26,208 + exit 0 = 348,040.

Raw stockout probability: D0 .919, D1 .008, D2 .895. Display: 92%, 1%, 90%. The exact D1-D0 delta is -91.1 percentage points, not the difference between rounded labels.

Demand -15% TCO: D0 310,547, D1 358,161, D2 306,401. Lowest cost alone does not establish feasibility. Lead-time stress: D0 361,502, D1 347,640, D2 356,665. See final-numbers.csv. Service level is not 1 minus stockout probability.

Reproducibility and performance

The 10,000-run local preview took approximately 0.51 seconds. Repeating a fixed seed produced the same matrix hash. Seeds 1042026, 1042027 and 1042028 differed by at most 0.78 percentage points in stockout probability. This does not replace the reviewed 1,000-run public result. Peak RSS was not measured. AI network latency is separate from pure simulation time.

Regression and handoff

Backend 24 tests, AI/evaluation 125 tests, frontend 62 tests and static deployment 3 tests passed. Four Golden Evidence records and tamper rejection passed. Fourteen offline failure cases passed. Automated checks do not constitute a participant's signature.

Wang Hao's historical feedback is provided as a labelled English summary. Deng's unperformed user tests are not reported as completed. Follow Recording_Script.md and Screenshot_Checklist.md. Keep LIVE/CACHED labels visible. Do not splice mock outputs into a live demonstration. The 160-second script is ready, but final recording and upload remain Day7 work.

Shared Release and Acceptance

Public release

Live: <https://causora-one.vercel.app/> . Repository: <https://github.com/Ding-hub6029/Causora> . Horizon 104 weeks / 24 months, ds-001, seed 1042026, tco-v1. Public Golden uses 1,000 runs. Local benchmark uses 10,000 runs.

Verification

149 combined backend/AI/evaluation tests passed. Three admission tests passed, including two already counted in 149. Frontend 62 and static deployment 3 passed. Type checking, lint and live/static builds passed. Golden has four validated Evidence records and rejects tampering. Fourteen offline failure fixtures passed.

Numbers and performance

Baseline D1 TCO is USD 348,040, USD 17,771 below D0. Raw stockout probability is .008, displayed as about 1%. All nine cost decompositions reconcile. Local 10,000-run computation took about 0.51 seconds. Same-seed hashes match. Three-seed stockout variation is at most 0.78 percentage points. Peak RSS was not measured.

Real AI and cache

Real provider receipts include generation IDs, output hashes and costs. Code computes numbers. Models generate qualitative reviews. No claim is made to inspect private model reasoning. Cached runs are historical, validated and read-only. They are not new calls.

Admission, budget and privacy

One review at a time in one worker, with a 15-second minimum start interval. Rejected 429 requests make no model call. PostgreSQL preserves budget history. Spending uses the existing teacher-key balance only. Each review permits six calls and zero transport retries. Credentials are excluded from delivery.

Acceptance still pending

The final 2-4 minute video, public upload, formal Devpost submission, member-account verification and current-round genuine incognito/unfamiliar-device acceptance remain pending. Only one real user, Wang Hao, is evidenced. Real AI under 45 seconds, ten online successes and independent unseen Critic recall $\geq 80\%$ are not established.

Day7 handoff

Record the real flow using the 160-second script. Upload a publicly playable video, verify member accounts and all public links, then submit. Fix bugs only. Formula, data, model or prompt changes require regression and a new Golden.