

Day6 Delivery and Verification

Ding Xiangfeng

Deployment and demonstration

The package includes deployable source, validated Golden, live/cache distinctions, G6 gates, a Devpost draft, recording script and recovery instructions.

1. Open the public URL and wait for Simulation service reachable. Health readiness alone does not prove AI completion.
2. Run simulation in Decision Matrix. Wait for LIVE API and all nine cells. Open D1 Formula Trace.
3. Run the matching Boardroom. A usable Brief requires a complete Critic, passed Numeric Guardrail and matching run identity.
4. On 429, wait for the current review. On 503 or timeout, retry Boardroom while retaining the Matrix. After a backend restart, generate a fresh simulation if the old identity is no longer valid.
5. During cold starts, use Verified cache and identify it as a validated historical real run, not a new AI call. Never present Day2 examples as live results.
6. Actually test a fresh incognito session and a teammate's mobile connection before signing acceptance. These checks are not fabricated here.

Recovery and release

Render root: Day5/Causora. Start: `python deployment/render_start.py`. Vercel uses the frontend live build and public API URL. Backend credentials and PostgreSQL settings must not enter NEXT_PUBLIC variables.

Preserve budget history. Do not recharge or buy subscriptions. Single-worker admission protection is enabled. Data remains synthetic and the README limitations must remain.

See G6_Gates.md. Add the actual uploaded video URL to Devpost.md. Real members must confirm accounts. Day7 covers recording, upload, link verification and submission. Do not claim submission is complete.

Shared Release and Acceptance

Public release

Live: <https://causora-one.vercel.app/> . Repository: <https://github.com/Ding-hub6029/Causora> . Horizon 104 weeks / 24 months, ds-001, seed 1042026, tco-v1. Public Golden uses 1,000 runs. Local benchmark uses 10,000 runs.

Verification

149 combined backend/AI/evaluation tests passed. Three admission tests passed, including two already counted in 149. Frontend 62 and static deployment 3 passed. Type checking, lint and live/static builds passed. Golden has four validated Evidence records and rejects tampering. Fourteen offline failure fixtures passed.

Numbers and performance

Baseline D1 TCO is USD 348,040, USD 17,771 below D0. Raw stockout probability is .008, displayed as about 1%. All nine cost decompositions reconcile. Local 10,000-run computation took about 0.51 seconds. Same-seed hashes match. Three-seed stockout variation is at most 0.78 percentage points. Peak RSS was not measured.

Real AI and cache

Real provider receipts include generation IDs, output hashes and costs. Code computes numbers. Models generate qualitative reviews. No claim is made to inspect private model reasoning. Cached runs are historical, validated and read-only. They are not new calls.

Admission, budget and privacy

One review at a time in one worker, with a 15-second minimum start interval. Rejected 429 requests make no model call. PostgreSQL preserves budget history. Spending uses the existing teacher-key balance only. Each review permits six calls and zero transport retries. Credentials are excluded from delivery.

Acceptance still pending

The final 2-4 minute video, public upload, formal Devpost submission, member-account verification and current-round genuine incognito/unfamiliar-device acceptance remain pending. Only one real user, Wang Hao, is evidenced. Real AI under 45 seconds, ten online successes and independent unseen Critic recall $\geq 80\%$ are not established.

Day7 handoff

Record the real flow using the 160-second script. Upload a publicly playable video, verify member accounts and all public links, then submit. Fix bugs only. Formula, data, model or prompt changes require regression and a new Golden.