

Day6 Delivery and Verification

Wang Hao

Provider and fallback

Real OpenRouter receipts and offline fixtures are archived separately. freeze.json records prompt-source hashes, models, budget policy, six calls per review, zero transport retries and run identity.

Primary: openai/gpt-5-mini. Preferred Critic: google/gemini-3.1-pro-preview. Validated same-family GPT fallback applies when the preferred Critic is unavailable. CFO, COO and Risk receive bounded inputs. Critic examines cross-functional omissions. Brief numerical claims reference code-computed values.

Each review permits at most six calls. A correction can consume the remaining call, with no unlimited automatic retry. Wrong numbers, citations, run identity, scenario or missing Critic are rejected. Failure preserves the validated Matrix and Formula Trace.

Fourteen offline cases cover missing/timed-out/malformed Critic, injection, wrong citations, numerical tampering, wrong roles, zero-call no-feasible handling, Matrix preservation and D2 semantics. A 14/14 fixture result is not real model recall.

Admission and budget

One public review can run at a time, with at least 15 seconds between starts. Rejection returns 429 and Retry-After. Exceptions release the slot. Health and simulation remain available. Admission is process-local and requires distributed protection before horizontal scaling.

PostgreSQL row locking protects persistent budget history across processes. There is no cumulative call-count cap, but spending is limited to the existing teacher-key balance and a frozen ceiling that does not grow with recharges. Per-review reservations remain. Reserved amounts are not actual charges.

Cache, privacy and future changes

Golden is a read-only historical real run checked for source integrity and matching identity. Tampering fails and replay does not bill a new call. Day2 examples cannot be approved. Only a fresh successful matching live Brief can record a browser-local choice.

Packages exclude API keys, database URLs, private environment files and credentials. Receipts retain generation IDs, output hashes, statuses and costs. Human review remains necessary because arbitrary semantic evidence entailment is not guaranteed.

Do not force D1 or force Critic to claim a finding. Follow actual data and constraints. Model or prompt changes require regression and a new freeze. Describe the system accurately: AI generates qualitative reviews while code computes and constrains numerical results.

Shared Release and Acceptance

Public release

Live: <https://causora-one.vercel.app/> . Repository: <https://github.com/Ding-hub6029/Causora> . Horizon 104 weeks / 24 months, ds-001, seed 1042026, tco-v1. Public Golden uses 1,000 runs. Local benchmark uses 10,000 runs.

Verification

149 combined backend/AI/evaluation tests passed. Three admission tests passed, including two already counted in 149. Frontend 62 and static deployment 3 passed. Type checking, lint and live/static builds passed. Golden has four validated Evidence records and rejects tampering. Fourteen offline failure fixtures passed.

Numbers and performance

Baseline D1 TCO is USD 348,040, USD 17,771 below D0. Raw stockout probability is .008, displayed as about 1%. All nine cost decompositions reconcile. Local 10,000-run computation took about 0.51 seconds. Same-seed hashes match. Three-seed stockout variation is at most 0.78 percentage points. Peak RSS was not measured.

Real AI and cache

Real provider receipts include generation IDs, output hashes and costs. Code computes numbers. Models generate qualitative reviews. No claim is made to inspect private model reasoning. Cached runs are historical, validated and read-only. They are not new calls.

Admission, budget and privacy

One review at a time in one worker, with a 15-second minimum start interval. Rejected 429 requests make no model call. PostgreSQL preserves budget history. Spending uses the existing teacher-key balance only. Each review permits six calls and zero transport retries. Credentials are excluded from delivery.

Acceptance still pending

The final 2-4 minute video, public upload, formal Devpost submission, member-account verification and current-round genuine incognito/unfamiliar-device acceptance remain pending. Only one real user, Wang Hao, is evidenced. Real AI under 45 seconds, ten online successes and independent unseen Critic recall $\geq 80\%$ are not established.

Day7 handoff

Record the real flow using the 160-second script. Upload a publicly playable video, verify member accounts and all public links, then submit. Fix bugs only. Formula, data, model or prompt changes require regression and a new Golden.